

# ECS 150 - Project 4

---

*Prof. Joël Porquet*

UC Davis - Spring Quarter 2019



# Goal

---

The goal of this project is to implement the support for a very simple *file system*: **ECS150-FS**. Applications will have the possibility to read/write files from/to this file system.

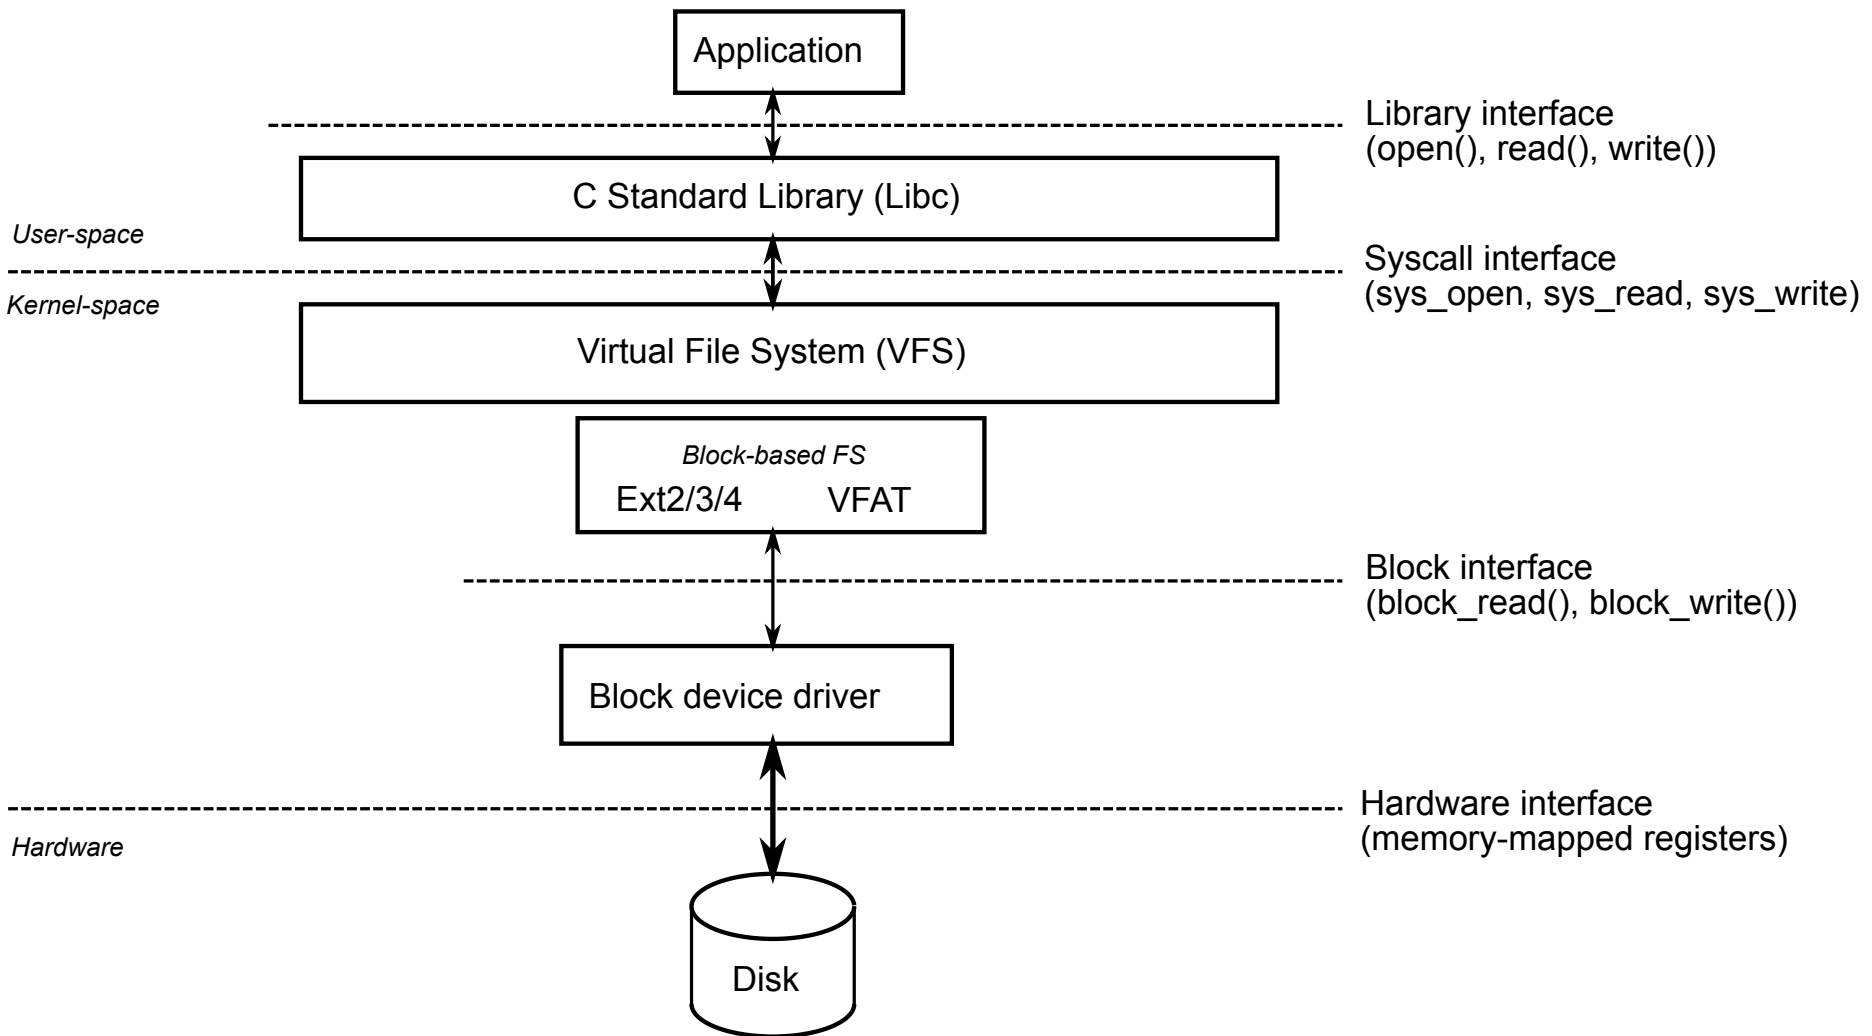
```
int fs_mount(const char *diskname);           /* Mounting the file system */
int fs_umount(void);
int fs_info(void);

int fs_create(const char *filename);          /* Creating files */
int fs_delete(const char *filename);
int fs_ls(void);

int fs_open(const char *filename);            /* Opening files */
int fs_close(int fd);
int fs_stat(int fd);
int fs_lseek(int fd, size_t offset);

int fs_write(int fd, void *buf, size_t count); /* Modifying files */
int fs_read(int fd, void *buf, size_t count);
```

# Big picture: reality



## Problems:

- The vast majority the file system management is in kernel mode!
- And we need a physical disk...

# Emulating a disk with a file

A disk, or a partition on a disk, merely represents contiguous binary data storage.

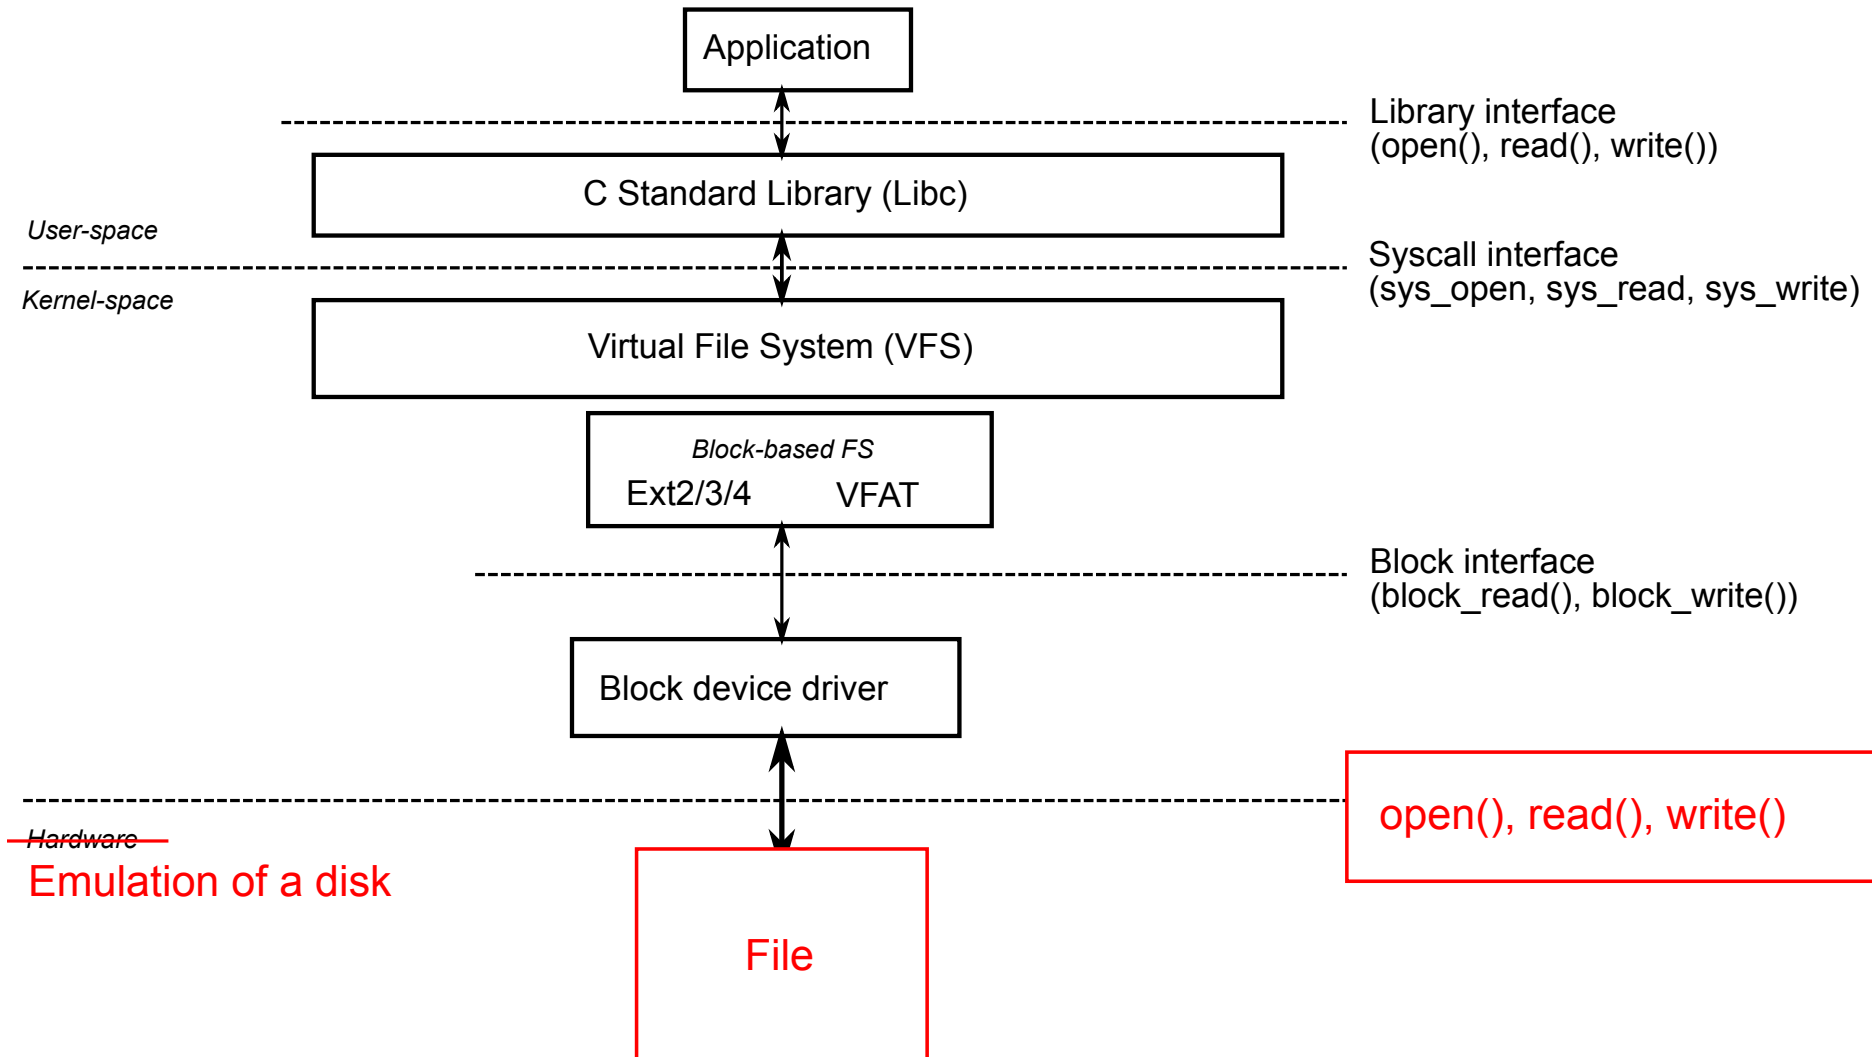
How can we easily emulate any size of contiguous data?... With a file!

```
$ dd if=/dev/zero of=emulated_disk_space bs=4K count=8192

$ ls -l emulated_disk_space
-rw-r--r-- 1 joel joel 32M 2017-03-01 13:52 emulated_disk_space

$ xxd emulated_disk_space
00000000: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000010: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000020: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000030: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000040: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000050: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000060: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000070: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000080: 0000 0000 0000 0000 0000 0000 0000 0000  .....
00000090: 0000 0000 0000 0000 0000 0000 0000 0000  .....
...
01fffff0: 0000 0000 0000 0000 0000 0000 0000 0000  .....
```

# Replacing the disk



# Accessing a file by blocks

---

```
#define BLOCK_SIZE 4096

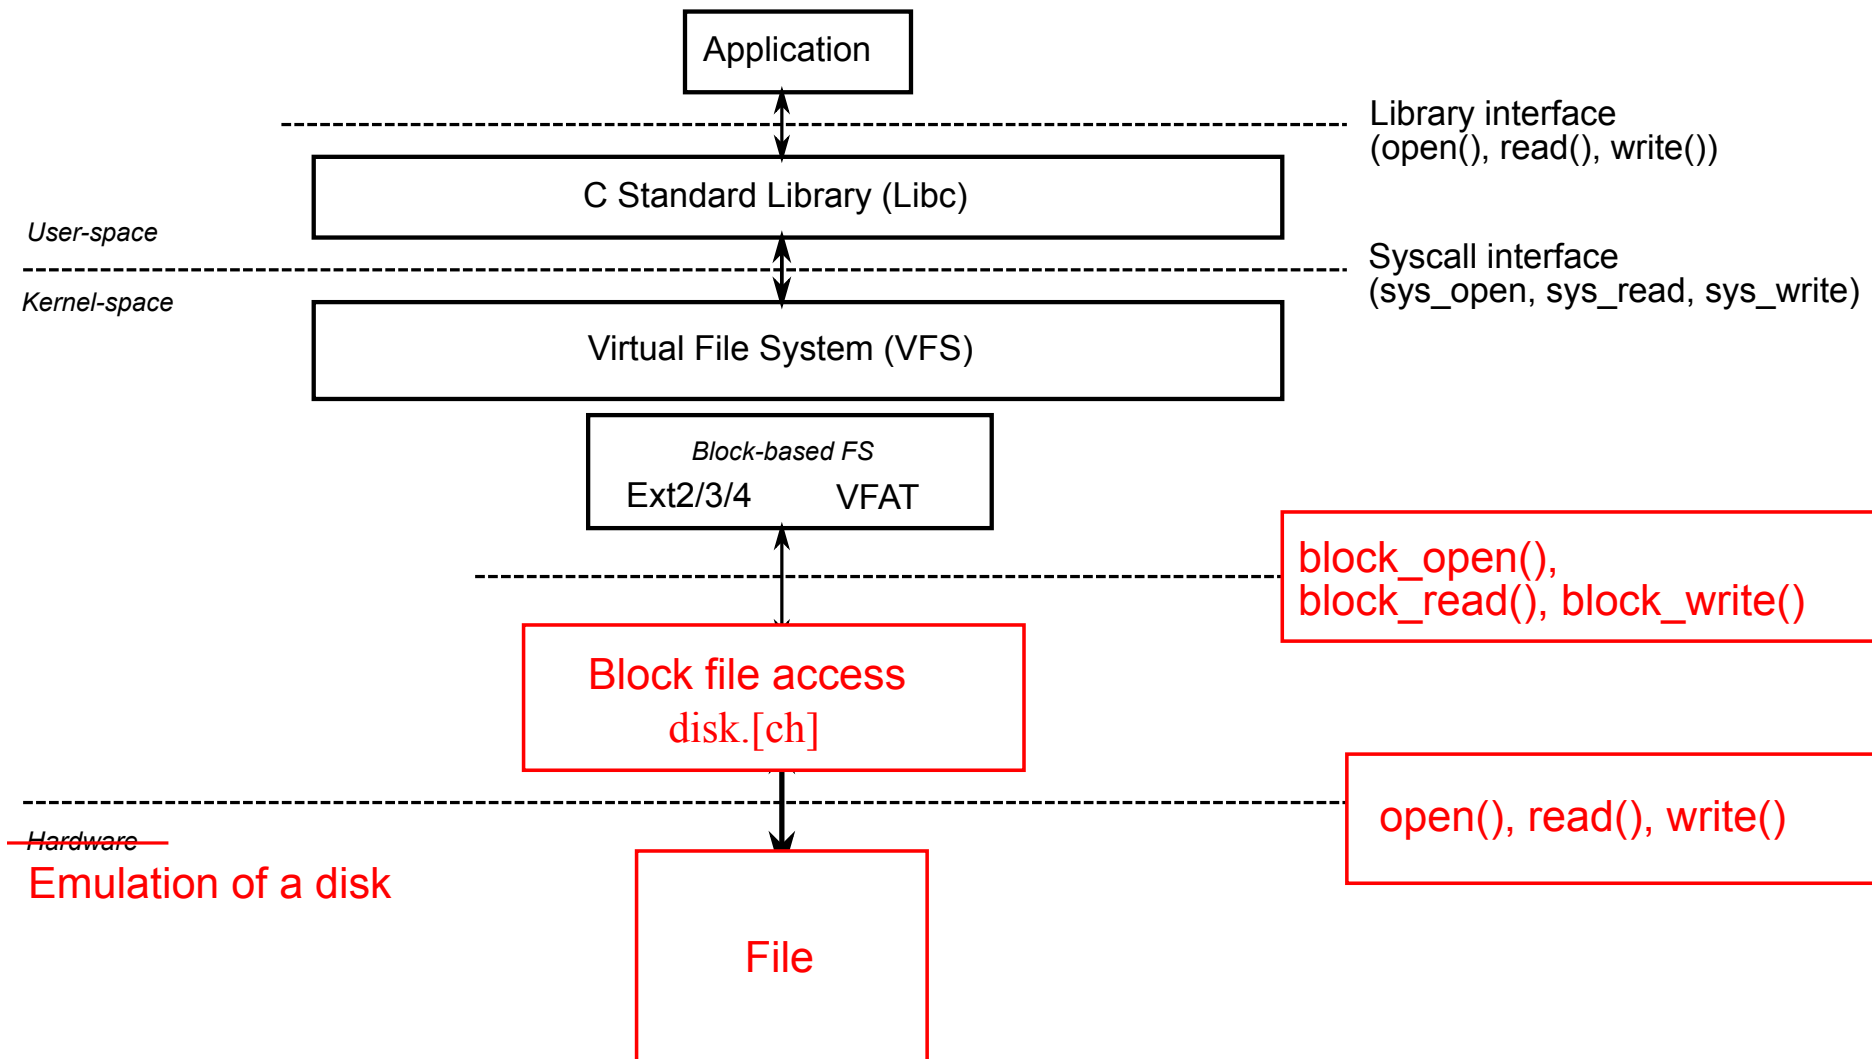
int fd;

int block_open(char *disk_filename)
{
    fd = open(disk_filename, O_RDWR);
}

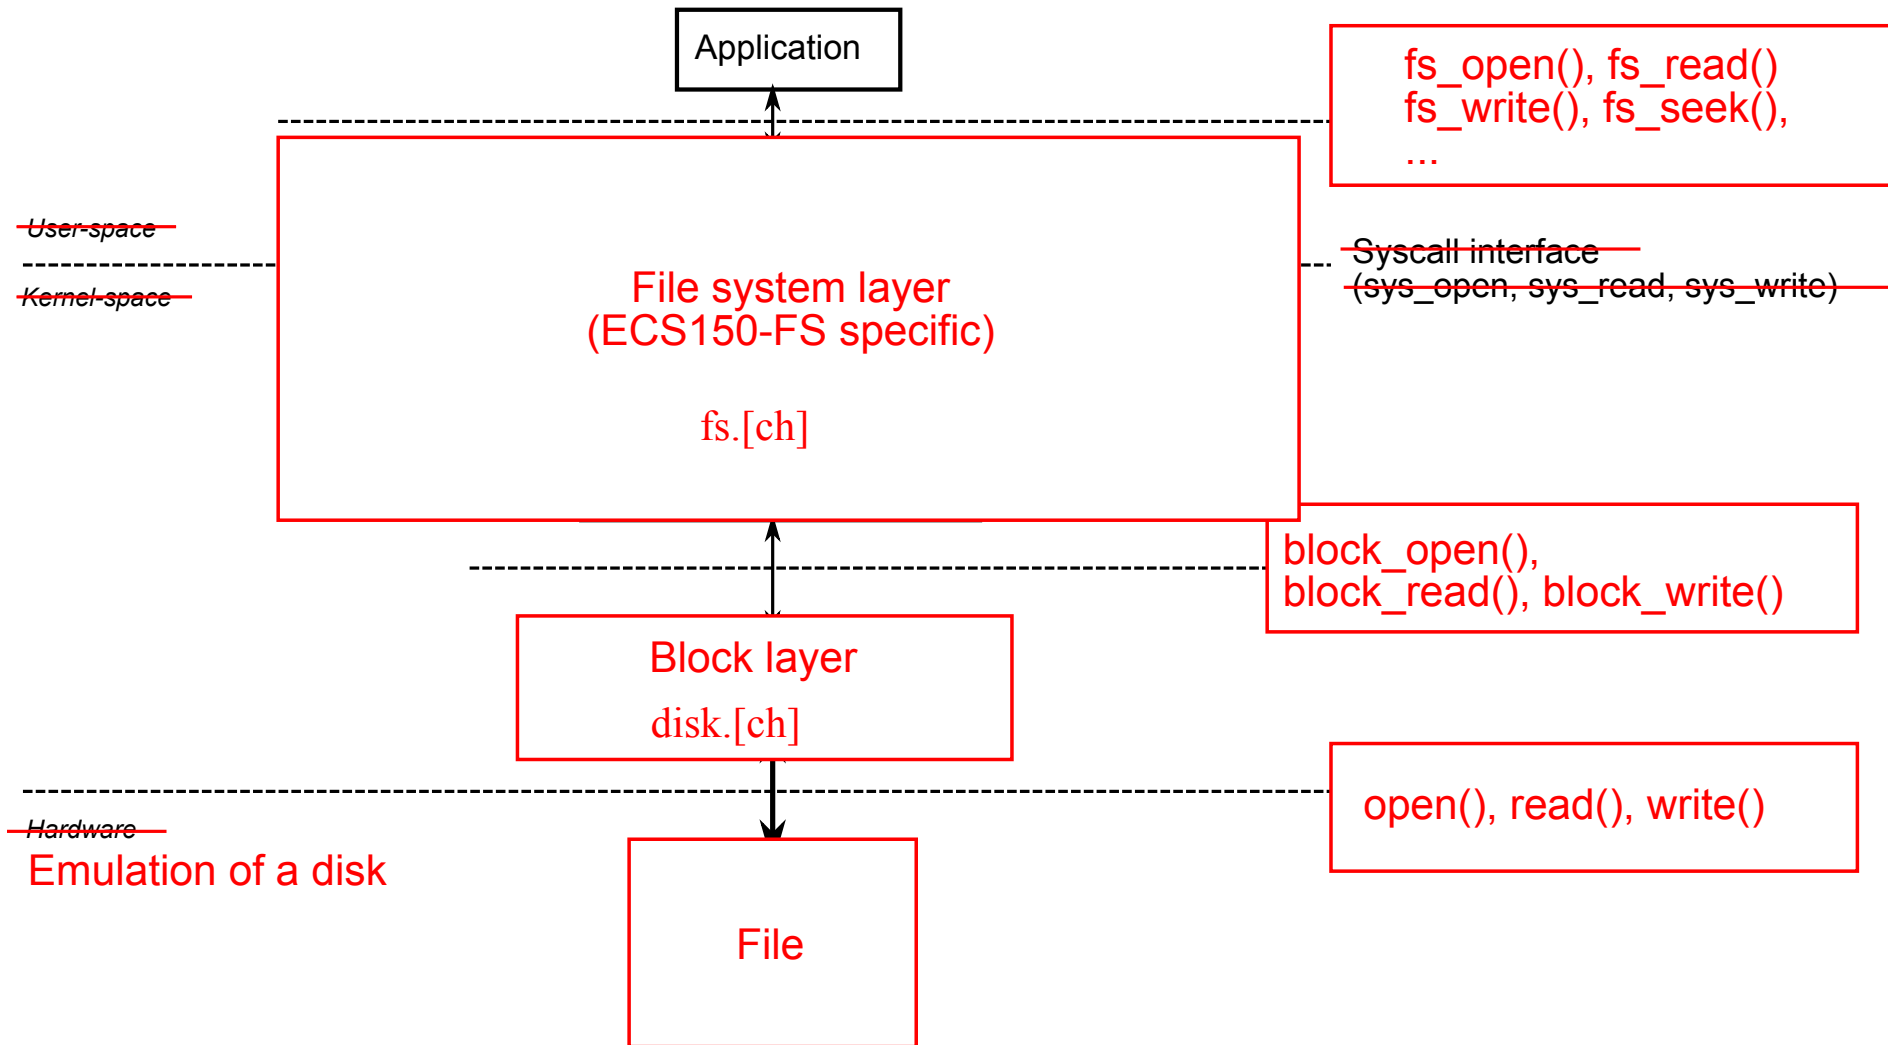
int block_read(size_t block_nr, void *buf)
{
    lseek(fd, block_nr * BLOCK_SIZE);
    read(fd, buf, BLOCK_SIZE);
}

int block_write(size_t block_nr, void *buf)
{
    lseek(fd, block_nr * BLOCK_SIZE);
    write(fd, buf, BLOCK_SIZE);
}
```

# Replacing the block device driver



# Replacing the libc/vfs/fs drivers



- *That's the project!*



# ECS150-FS

## Layout



Each block is 4096 bytes.

Example with file system embedding 8192 data blocks:

```
$ ./fs_make.x disk.fs 8192  
Creating virtual disk 'disk.fs' with '8192' data blocks
```

- Amount of data blocks: 8192
- Number of blocks for FAT:  $(8192 * 2) / 4096 = 4$
- Total amount of blocks:  $1 + 4 + 1 + 8192 = 8198$
- Root directory block index: 5
- Data block start index: 6


# ECS150-FS

## Superblock: high-level layout

| Offset | Length (bytes) | Description                             |
|--------|----------------|---|
| 0x00   | 8              | Signature (must be equal to "ECS150FS") |
| 0x08   | 2              | Total amount of blocks of virtual disk  |
| 0x0A   | 2              | Root directory block index              |
| 0x0C   | 2              | Data block start index                  |
| 0x0E   | 2              | Amount of data blocks                   |
| 0x10   | 1              | Number of blocks for FAT                |
| 0x11   | 4079           | Unused/Padding                          |

# ECS150-FS

## Superblock: at byte level



# ECS150-FS

---

## Superblock: C data structure

```
struct superblock{  
    ???  
};
```

Key points:

- The integer types must match exactly those of the specification
- It must represent the entirety of the block, padding included!

# Digression

## Integer types

- Is `char` always 8 bits?
- Is `short int` always 16 bits?
- Is `int` always 32 bits?
- Etc.

| Type  | Specification  |
|-------|--|
| char  | "Smallest addressable unit of the machine that can contain basic character set"  |
| short | "Capable of containing <i>at least</i> the $[-32767, +32767]$ range; thus, it is <i>at least</i> 16 bits in size."           |
| int   | "Capable of containing <i>at least</i> the $[-32767, +32767]$ range; thus, it is <i>at least</i> 16 bits in size."           |
| long  | "Capable of containing <i>at least</i> the $[-2147483647, +2147483647]$ range; thus, it is <i>at least</i> 32 bits in size." |

How to guarantee a certain size then?

# Digression

---

## Integer types

Use integer types that have exact widths (since C99):

```
#include <stdint.h>
```

```
int8_t  
int16_t  
int32_t
```

```
uint8_t  
uint16_t  
uint32_t
```

# Digression

## Reading data structures from a file (or buffer)

Reading data from a file (or whatever blob of data), for which I know the layout. It can easily be type-casted into a structure instance.


```
struct mystruct
{
    int32_t    a;
    int16_t    b;
    int16_t    c;
    int32_t    padding[2];
};
```

```
fd = open("file", O_RDWR);
...
char* buf[16];
read(fd, buf, 16);

struct mystruct *s = buf;

printf("s->a = %d\n", s->a);
s->a = 0;

write(fd, buf, 16);
```



```
/* or simply */
struct mystruct obj;
read(fd, &obj, sizeof(obj));

printf("obj.a = %d\n", obj.a);
obj.a = 0;

write(fd, &obj, sizeof(obj));
```

# ECS150-FS

## Layout




## FAT

- Big array of 16-bit entries: linked-list of data blocks composing a file
- Three possible values for each entry:
  - 0: corresponding data block is available
  - FAT\_EOC: last data block of a file
  - !=0 && !=FAT\_EOC: index of next data block



# ECS150-FS

## Layout



## Root directory

1 block, 32-byte entry per file: 128 entries total

| Offset | Length (bytes) | Description                         |
|--------|----------------|-------------------------------------|
| 0x00   | 16             | Filename (including NULL character) |
| 0x10   | 4              | Size of the file (in bytes)         |
| 0x14   | 2              | Index of the first data block       |
| 0x16   | 10             | Unused/Padding                      |

# ECS150-FS

Example: big file, small file, empty file

Root directory

<"test1", 22000, 2>,  
<"test2", 5000, 1>,  
<"test3", 0, FAT\_EOC>  
...

FAT

|    |         |
|----|---------|
| 0  |         |
| 1  | 4       |
| 2  | 3       |
| 3  | 5       |
| 4  | FAT_EOC |
| 5  | 6       |
| 6  | 7       |
| 7  | 8       |
| 8  | FAT_EOC |
| 9  | 0       |
| 10 | 0       |

Data blocks

|    |                   |
|----|-------------------|
| 0  |                   |
| 1  | (test2, block #0) |
| 2  | (test1, block #0) |
| 3  | (test1, block #1) |
| 4  | (test2, block #1) |
| 5  | (test1, block #2) |
| 6  | (test1, block #3) |
| 7  | (test1, block #4) |
| 8  | (test1, block #5) |
| 9  |                   |
| 10 |                   |

# Implementation

---

## Phase 1: Volume mounting

- `fs_mount()`:
  - Open the virtual disk
  - Read the metadata (superblock, fat, root directory)
- `fs_unmount()`:
  - Close virtual disk (make sure that virtual disk is up-to-date)
- `fs_info()`:
  - Show information about volume

# Implementation

---

## Phase 2: File creation/deletion

- `fs_create()`:
  - Create a new file
  - Initially, size is 0 and pointer to first data block is `FAT_EOC`
- `fs_delete()`:
  - Delete an existing file
  - Don't forget to free allocated data blocks
- `fs_ls()`:
  - List all the existing files

# Implementation

---

## Phase 3: File descriptor operations

- `fs_open()`:
  - Initialize and return file descriptor
  - 32 file descriptors max
  - Can open same file multiple times
  - Contains file's offset (initially 0)
- `fs_close()`:
  - Close file descriptor
- `fs_seek()`:
  - Move file's offset
- `fs_stat()`:
  - Return file's size

*None of these functions should change the file system...*

# Implementation

---

## Phase 4: File reading/writing

Most complicated phase: might take as much time as all the previous phases combined

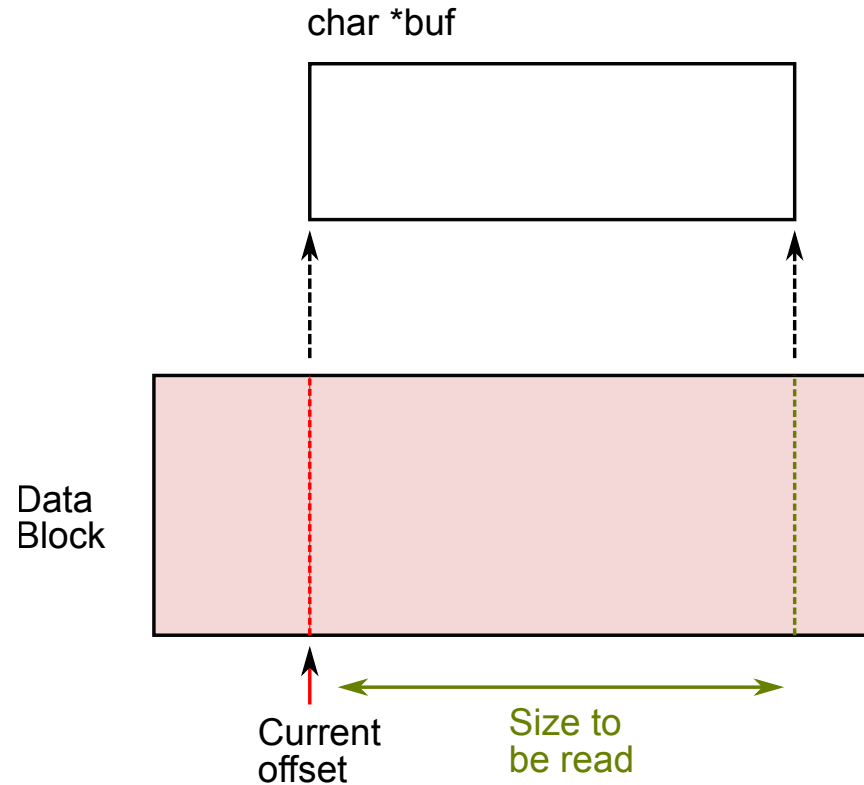
Three difficulties:

1. Small operations
2. First/last block on big operations
3. Extending writes

# Implementation

## Small operation: example

- Current offset is in the middle of the file, not aligned on the beginning of a block
- The size of data to read is smaller than what's remaining in this block




Might want to use a *bounce buffer*

# Implementation

## Big operation: example

- Current offset is in the middle of the file, not aligned on the beginning of a block
- The size to read spans multiple (non-consecutive) blocks
- The size of data to read is smaller than what's remaining in the last block




Mix of *bounce buffer* and direct copy



# Implementation

## Extending write: example

- Write more than what's currently allocated



- Allocation of new blocks must follow *first-fit* strategy (allocate first free data block from beginning of the FAT).

# Implementation

---

## In short

- Think of all the cases: combination of file's offset, file's size, size to be read or written, etc.
- Come up with a way to handle all these combinations in the *most* generic way (i.e., not one function or one separate if statement per case!)